

Adaptive Pareto Smoothed Importance Sampling

William Ruth

Joint work with Payman Nickchi

RichCon!

Importance Sampling

- Need to compute an expected value
 - $\mathbb{E}_F \varphi(X) = \int \varphi(x) f(x) dx$
- Can't do the integral
- Monte Carlo approximation:
 - Simulate from F
 - Average over simulations
- $\hat{\mathbb{E}} = \sum_i \frac{\varphi(X_i)}{M}, X_i \stackrel{\text{iid}}{\sim} F$

Importance Sampling

- Simulating from F might be hard
- “Multiply by 1”:

$$\begin{aligned}\mathbb{E}_F \varphi(X) &= \int \varphi(x) f(x) dx \\ &= \int \varphi(x) \frac{f(x)}{g(x)} g(x) dx \\ &= \mathbb{E}_G \left[\varphi(X) \cdot \frac{f(X)}{g(X)} \right] \\ &= \mathbb{E}_G [\varphi(X) \cdot w(X)]\end{aligned}$$

Importance Sampling

$$\mathbb{E}_F \varphi(X) = \mathbb{E}_G [\varphi(X) \cdot w(X)]$$

- G can be anything*
 - *Some choices will be better than others
- Simulate from G to estimate $\mathbb{E}_G [\varphi(X) \cdot w(X)]$
 - By extension, estimate $\mathbb{E}_F \varphi(X)$

$$\hat{\mathbb{E}} = \sum_i \frac{\varphi(X_i) \cdot w(X_i)}{M}, X_i \stackrel{\text{iid}}{\sim} G$$

Example: Mystery Target

- f unknown, but can be evaluated
- Try some proposals:
 - $G_1 \sim N(0, 1)$
 - $G_2 \sim N(2, 1)$
- Use $M = 1000$ samples from proposal
 - $\hat{\mathbb{E}}_1 =$
 - $\hat{\mathbb{E}}_2 =$

Example: Mystery Target

Histograms of weights

Importance Sampling

- Can we quantify this difference?
 - Yes!
- “Effective Sample Size”

$$ESS = \frac{M}{\sum_i w(X_i)^2} \leq M$$

Example: Mystery Target

Histograms of weights with ESS

Importance Sampling

- Problem: Low ESS \rightarrow can't estimate means
- But ESS *is* a mean
 - (Chatterjee and Diaconis, 2018)

Improving IS

- Large discrepancies in weights is bad
 - Reduce discrepancy
 - Shrink large weights
- Truncated IS
- Pareto Smoothed IS

Improving IS

- Truncated IS (Ionides, 2008):
 1. Choose a threshold
 2. Set any weights above threshold equal to threshold

Example: Mystery Target

Histograms of weights with threshold

Example: Mystery Target

Histograms of truncated weights

Example: Mystery Target

Histograms of truncated weights with before and after ESS

Improving IS

- Pareto Smoothing (Vehtari et al., 2022):
 1. Choose a threshold
 - Weights above threshold represent tail of their dist.
 2. Approximate tail with Generalized Pareto Dist.
 - Fit GPD to weights above threshold
 3. Replace large weights with quantiles of fitted GPD

Example: Mystery Target

Histograms of weights with threshold

Example: Mystery Target

Histograms of weights with threshold and fitted GPD density above threshold

Example: Mystery Target

Histograms of smoothed weights

Example: Mystery Target

Histograms of smoothed weights with ESS for raw, truncated and smoothed weights

Adaptive IS

- Modifications are nice, but require creativity
- Alternative: directly optimize ESS
- Adaptive Importance Sampling (Akyildiz and Míguez, 2021)
 - Choose a family of proposals
 - Iteratively update the proposal to maximize ESS

Adaptive IS

- Recall:

$$ESS = \frac{M}{\sum_i w(X_i)^2} =: \frac{M}{\hat{\rho}}$$

- Want to maximize a population-level analog
 - Equivalently, minimize $\rho = \mathbb{E}_G [w(X)^2]$
- We only get ESS, $\hat{\rho}$
- Noisy version of the function we want to optimize

Adaptive IS

- Stochastic Approximation:
- If we had ρ , do gradient descent
- $\theta_{k+1} = \theta_k - \alpha \nabla \rho(\theta_k)$
- Instead, do gradient descent on $\hat{\rho}$
- $\hat{\theta}_{k+1} = \hat{\theta}_k - \alpha_k \nabla \hat{\rho}(\hat{\theta}_k)$

Adaptive IS

- Overview (cite Akyldiz and Miguez)
- Stochastic approximation

Pareto Tail Diagnostic

- Alt. to ESS
- Mention Chattergee and Diaconis

Adaptive Pareto Smoothed IS

- Discuss what we're doing
- Gaussian example?

Acknowledgements



Thank You

Some References

- Akyildiz, Ö. D. and Míguez, J. (2021). Convergence rates for optimized adaptive importance samplers. *Statistics and Computing*, 31(12).
- Chatterjee, S. and Diaconis, P. (2018). The sample size required in importance sampling. *The Annals of Applied Probability*, 28(2).
- Ionides, E. L. (2008). Truncated importance sampling. *Journal of Computational and Graphical Statistics*, 17(2).
- Vehtari, A., Simpson, D., Gelman, A., Yao, Y., and Gabry, J. (2022). Pareto smoothed importance sampling. *ArXiv*.